

COMP4680/8650: Advanced Topics in Machine Learning

Weeks 9–10 — Deep Learning

Stephen Gould

`stephen.gould@anu.edu.au`

Australian National University

Semester 2, 2022

Optimization and Machine Learning

Generic mathematical optimization problems have the form

$$\begin{array}{ll}\text{minimize} & f_0(x) \\ \text{subject to} & f_i(x) \leq b_i, \quad i = 1, \dots, m\end{array}$$

The solution $x^* = (x_1^*, \dots, x_n^*)$ can be:

- ▶ investment amounts in portfolio optimization
- ▶ device sizes in electronic circuits
- ▶ equipment settings in manufacturing
- ▶ **model parameters θ in machine learning**, e.g.,
 - ▶ probability distribution, $p_\theta : \mathcal{A} \rightarrow [0, 1]$
 - ▶ prediction function, $h_\theta : \mathcal{X} \rightarrow \mathcal{Y}$

Deep learning is a sub-field of machine learning focusing on *end-to-end* learnable models based on artificial neural networks.

Multilayer Perceptron (MLP)

- ▶ most basic feedforward artificial neural network, originating from work of Rosenblatt (1961) and Widrow & Hoff (1960)
- ▶ input $x \in \mathbb{R}^{p_0}$, hidden layers $z_i \in \mathbb{R}^{p_i}$, output $y \in \mathbb{R}^{p_n}$
- ▶ each layer computes its output as

$$z_i = \tilde{f}_i(z_{i-1}) = f_i(A_i z_{i-1} + b_i)$$

where $A_i \in \mathbb{R}^{p_i \times p_{i-1}}$ and $b_i \in \mathbb{R}^{p_i}$ are parameters and f_i is a non-linear (elementwise) activation function ($x \triangleq z_0$, $y \triangleq z_n$)

- ▶ the network output is then the composition

$$\begin{aligned} y &= (\tilde{f}_n \circ \cdots \circ \tilde{f}_2 \circ \tilde{f}_1)(x) \\ &= f_n(A_n f_{n-1}(\cdots f_1(A_1 x + b_1)) + b_n) \end{aligned}$$

Activation Functions

- ▶ logistic function (sigmoid)


$$f(x) = \frac{1}{1 + e^{-x}}$$

- ▶ hyperbolic tangent

$$f(x) = \tanh(x) = \frac{e^{2x} - 1}{e^{2x} + 1}$$


- ▶ rectified linear unit (ReLU)

$$f(x) = \max\{0, x\}$$



Single Neuron

- a single neuron computes $y = f(\sum_{i=1}^n a_i x_i + b) \in \mathbb{R}$



(Supervised) Learning Review

- ▶ parametrized model $y = h_{\theta}(x)$ and training set $\mathcal{D} = \{(x_i, y_i)\}$
- ▶ regularized loss function \mathcal{L} that measures how well the model fits the training data, usually decomposes over the training data plus a prior over parameters

$$\mathcal{L}(\theta) = \sum_i \ell(h_{\theta}(x_i), y_i) + R(\theta)$$

- ▶ we then optimize the loss function with respect to the model parameters

$$\theta^{\star} = \underset{\theta}{\operatorname{argmin}} \mathcal{L}$$

- ▶ simplest optimization algorithm is then gradient descent

$$\theta \leftarrow \theta - \eta \nabla_{\theta} \mathcal{L}$$

Back-propagation

- ▶ consider a single neuron $y = f(a^T x + b)$
- ▶ if we have $\frac{\partial \mathcal{L}}{\partial y}$ we can compute $\frac{\partial \mathcal{L}}{\partial a_i}$, by the **chain rule**, as

$$\begin{aligned}\frac{\partial \mathcal{L}}{\partial a_i} &= \frac{\partial \mathcal{L}}{\partial y} \frac{\partial y}{\partial a_i} \\ &= \frac{\partial \mathcal{L}}{\partial y} f'(a^T x + b) x_i\end{aligned}$$

- ▶ likewise we can compute $\frac{\partial \mathcal{L}}{\partial b}$ and $\frac{\partial \mathcal{L}}{\partial x_i}$

Back-propagation: Why compute $\frac{\partial \mathcal{L}}{\partial z_i}$?

- ▶ now consider a composition of neurons

$$z_i = f_i(A_i z_{i-1} + b_i), \quad i = 1, \dots, n$$


- ▶ we care about updating all parameters of the network

$$\{(A_i, b_i) \mid i = 1, \dots, n\}$$

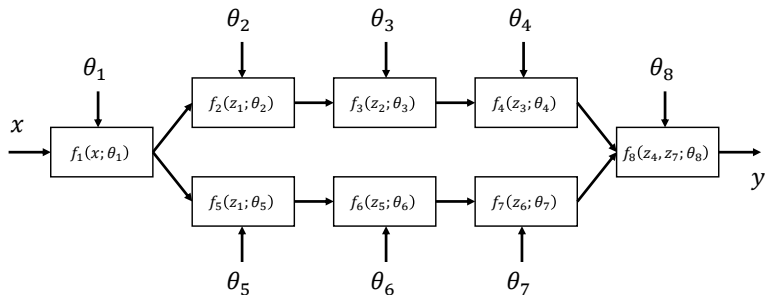
- ▶ if we have $\nabla_{z_i} \mathcal{L}$ we can update A_i and b_i
- ▶ to update A_{i-1} and b_{i-1} we need $\nabla_{z_{i-1}} \mathcal{L}$
- ▶ so computing $\frac{\partial \mathcal{L}}{\partial z_i}$ allows us to propagate gradients backwards through the network

Abstract Processing Node/Layer


- ▶ a neural network is a computation graph of processing nodes
- ▶ data (or gradient) tensors pass along edges between nodes



End-to-end Computation Graph




End-to-end Computation Graph



$$y = f_8(f_4(f_3(f_2(f_1(x))))), f_7(f_6(f_5(f_1(x))))))$$

(omitting parameters θ_i)

End-to-end Computation Graph Gradients



Example 1.

$$\frac{\partial \mathcal{L}}{\partial \theta_7} = \frac{\partial \mathcal{L}}{\partial y} \frac{\partial y}{\partial z_7} \frac{\partial z_7}{\partial \theta_7}$$

Example 2.

$$\frac{\partial \mathcal{L}}{\partial \theta_1} = \frac{\partial \mathcal{L}}{\partial y} \left(\frac{\partial y}{\partial z_4} \frac{\partial z_4}{\partial z_3} \frac{\partial z_3}{\partial z_2} \frac{\partial z_2}{\partial z_1} + \frac{\partial y}{\partial z_7} \frac{\partial z_7}{\partial z_6} \frac{\partial z_6}{\partial z_5} \frac{\partial z_5}{\partial z_1} \right) \frac{\partial z_1}{\partial \theta_1}$$

Back-propagation: MLP Worked Example

- ▶ two-layer perceptron with sigmoid activation and squared-loss

$$z = \sigma(Ax + b)$$

$$y = \sigma(c^T z + d)$$

$$\mathcal{L}(A, b, c, d) = \frac{1}{2}(y - t)^2$$

where $\sigma(\xi) = (1 + e^{-\xi})^{-1}$ is applied elementwise

- ▶ using $\frac{\partial \mathcal{L}}{\partial y} = y - t$ and $\sigma'(\xi) = \sigma(\xi)(1 - \sigma(\xi))$ we can compute updates as ...

Back-propagation: MLP Worked Example

- ▶ layer 2 parameters:

$$\frac{\partial \mathcal{L}}{\partial d} = (y - t)y(1 - y)$$
$$\nabla_c \mathcal{L} = (y - t)y(1 - y)z$$

- ▶ layer 2 input / layer 1 output:


$$\nabla_z \mathcal{L} = (y - t)y(1 - y)c$$

- ▶ layer 1 parameters:

$$\nabla_b \mathcal{L} = \nabla_z \mathcal{L} \circ z \circ (1 - z)$$
$$\nabla_A \mathcal{L} = (y - t)y(1 - y) (c \circ z \circ (1 - z)) x^T$$

- ▶ note that we have reused calculations from the forward pass (memory-vs-compute trade-off)

Memory Buffers



- ▶ in-place operations may save memory in the forward pass
- ▶ re-using buffers may save memory in the backward pass
- ▶ at *test time* buffering data is not needed

Automatic Differentiation

- ▶ automatic differentiation (AD) is an algorithmic procedure that produces code for computing the derivative of a function
- ▶ assumes numeric computations are composed of a finite set of elementary operations that we know how to differentiate
 - ▶ arithmetic, exp, log, trigonometric
- ▶ two flavours
 - ▶ (forward mode) propagate calculations of first-order approximation $x + \Delta x$ forward through the computations
 - ▶ (reverse mode) builds program to compute gradient based on the chain rule re-using computation where applicable
- ▶ different deep learning frameworks use slightly different approaches (explicit graph construction versus implicit operator tracking)
- ▶ python package **autograd** (for reverse mode AD)

Autograd Example

```
import autograd.numpy as np
from autograd import grad

# define a function
def tanh(x):
    y = np.exp(2.0 * x)
    return (y - 1.0) / (y + 1.0)

# create function to compute gradient
dtanh = grad(tanh)

# evaluate function and gradient at 1.0
print(tanh(1.0))
print(dtanh(1.0))
print((tanh(1.01) - tanh(0.99)) / 0.02)
```

Universal Approximation Theorem

(Cybenko 1989, Hornik 1991)

Theorem (roughly):

- ▶ let $f : \mathbb{R} \rightarrow \mathbb{R}$ be a (well-behaved) activation function
- ▶ let $g : \mathbb{R}^n \rightarrow \mathbb{R}$ be any continuous function
- ▶ then there exists parameters $A \in \mathbb{R}^{m \times n}$, $b \in \mathbb{R}^m$, and $c \in \mathbb{R}^m$ such that the function

$$\hat{g}(x) = c^T f(Ax + b)$$

approximates $g(x)$ everywhere

- ▶ that is, $|\hat{g}(x) - g(x)| \leq \epsilon$ for all x

Historical Notes

- ▶ neural networks were very popular from 1960s to 1990s
- ▶ non-convexity and problems with diminishing and exploding gradients made them difficult to train
- ▶ they also often had poor generalization to unseen data
- ▶ lack of theory compared to other machine learning approaches
- ▶ ... led to the AI winter

Historical Notes

- ▶ neural networks were very popular from 1960s to 1990s
- ▶ non-convexity and problems with diminishing and exploding gradients made them difficult to train
- ▶ they also often had poor generalization to unseen data
- ▶ lack of theory compared to other machine learning approaches
- ▶ ... led to the AI winter
- ▶ some persisted (Hinton, Lecun, Bengio, Schmidhuber, ...)
- ▶ Krizhevsky et al. 2012 started a resurgence in NN research
- ▶ we're now in the AI spring ...
- ▶ better training algorithms and software libraries
 - ▶ pre-training
 - ▶ stochastic gradient descent on mini-batches
 - ▶ drop-out (and other regularization incl. data augmentation)
 - ▶ batch normalization
- ▶ more labelled training data
- ▶ faster and bigger hardware (e.g., GPUs)
- ▶ still lack of theory but large community sharing best practice

Stochastic Gradient Descent (SGD)

- ▶ (batch) gradient descent on training set $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^m$ is

$$\begin{aligned}\theta &\leftarrow \theta - \eta \nabla_{\theta} \mathcal{L} \\ &= \theta - \eta \sum_{i=1}^m \nabla_{\theta} \ell(f(x_i; \theta), y_i)\end{aligned}$$

Stochastic Gradient Descent (SGD)

- ▶ (batch) gradient descent on training set $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^m$ is

$$\begin{aligned}\theta &\leftarrow \theta - \eta \nabla_{\theta} \mathcal{L} \\ &= \theta - \eta \sum_{i=1}^m \nabla_{\theta} \ell(f(x_i; \theta), y_i)\end{aligned}$$

- ▶ to save computation stochastic gradient descent approximates the gradient on mini-batches $\mathcal{I} \subseteq \{1, \dots, m\}$

$$\widehat{\nabla_{\theta} \mathcal{L}} = \sum_{i \in \mathcal{I}} \nabla_{\theta} \ell(f(x_i; \theta), y_i)$$

Stochastic Gradient Descent (SGD)

- ▶ (batch) gradient descent on training set $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^m$ is

$$\begin{aligned}\theta &\leftarrow \theta - \eta \nabla_{\theta} \mathcal{L} \\ &= \theta - \eta \sum_{i=1}^m \nabla_{\theta} \ell(f(x_i; \theta), y_i)\end{aligned}$$

- ▶ to save computation stochastic gradient descent approximates the gradient on mini-batches $\mathcal{I} \subseteq \{1, \dots, m\}$

$$\widehat{\nabla_{\theta} \mathcal{L}} = \sum_{i \in \mathcal{I}} \nabla_{\theta} \ell(f(x_i; \theta), y_i)$$

- ▶ can be accumulated over multiple forward passes

Stochastic Gradient Descent (SGD)

- ▶ (batch) gradient descent on training set $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^m$ is

$$\begin{aligned}\theta &\leftarrow \theta - \eta \nabla_{\theta} \mathcal{L} \\ &= \theta - \eta \sum_{i=1}^m \nabla_{\theta} \ell(f(x_i; \theta), y_i)\end{aligned}$$

- ▶ to save computation stochastic gradient descent approximates the gradient on mini-batches $\mathcal{I} \subseteq \{1, \dots, m\}$

$$\widehat{\nabla_{\theta} \mathcal{L}} = \sum_{i \in \mathcal{I}} \nabla_{\theta} \ell(f(x_i; \theta), y_i)$$

- ▶ can be accumulated over multiple forward passes
- ▶ under mild assumptions $E \left[\widehat{\nabla_{\theta} \mathcal{L}} \right] = \nabla_{\theta} \mathcal{L}$

Stochastic Gradient Descent (SGD)

- ▶ (batch) gradient descent on training set $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^m$ is

$$\begin{aligned}\theta &\leftarrow \theta - \eta \nabla_{\theta} \mathcal{L} \\ &= \theta - \eta \sum_{i=1}^m \nabla_{\theta} \ell(f(x_i; \theta), y_i)\end{aligned}$$

- ▶ to save computation stochastic gradient descent approximates the gradient on mini-batches $\mathcal{I} \subseteq \{1, \dots, m\}$

$$\widehat{\nabla_{\theta} \mathcal{L}} = \sum_{i \in \mathcal{I}} \nabla_{\theta} \ell(f(x_i; \theta), y_i)$$

- ▶ can be accumulated over multiple forward passes
- ▶ under mild assumptions $E \left[\widehat{\nabla_{\theta} \mathcal{L}} \right] = \nabla_{\theta} \mathcal{L}$
- ▶ in practice we permute $[m]$ and iterate through adjacent fixed-length intervals; once through the data is called an epoch

AdaGrad (Duchi et al., 2011)

- ▶ view gradient descent as a first-order proximal method

$$\begin{aligned}\theta^{(k+1)} &= \operatorname{argmin}_{\theta} \left\{ \langle \nabla \mathcal{L}(\theta^{(k)}), \theta \rangle + \frac{1}{2\eta_k} \|\theta - \theta^{(k)}\|_2^2 \right\} \\ &= \theta^{(k)} - \eta_k \nabla \mathcal{L}(\theta^{(k)})\end{aligned}$$

AdaGrad (Duchi et al., 2011)

- ▶ view gradient descent as a first-order proximal method

$$\begin{aligned}\theta^{(k+1)} &= \operatorname{argmin}_{\theta} \left\{ \langle \nabla \mathcal{L}(\theta^{(k)}), \theta \rangle + \frac{1}{2\eta_k} \|\theta - \theta^{(k)}\|_2^2 \right\} \\ &= \theta^{(k)} - \eta_k \nabla \mathcal{L}(\theta^{(k)})\end{aligned}$$

- ▶ adapt the step size (geometry) for different features to increase influence of rare but informative features

$$\begin{aligned}\theta^{(k+1)} &= \operatorname{argmin}_{\theta} \left\{ \langle \nabla \mathcal{L}(\theta^{(k)}), \theta \rangle + \frac{1}{2\eta_k} \|\theta - \theta^{(k)}\|_B^2 \right\} \\ &= \theta^{(k)} - \eta_k B^{-1} \nabla \mathcal{L}(\theta^{(k)})\end{aligned}$$

AdaGrad (Duchi et al., 2011)

- ▶ view gradient descent as a first-order proximal method

$$\begin{aligned}\theta^{(k+1)} &= \operatorname{argmin}_{\theta} \left\{ \langle \nabla \mathcal{L}(\theta^{(k)}), \theta \rangle + \frac{1}{2\eta_k} \|\theta - \theta^{(k)}\|_2^2 \right\} \\ &= \theta^{(k)} - \eta_k \nabla \mathcal{L}(\theta^{(k)})\end{aligned}$$

- ▶ adapt the step size (geometry) for different features to increase influence of rare but informative features

$$\begin{aligned}\theta^{(k+1)} &= \operatorname{argmin}_{\theta} \left\{ \langle \nabla \mathcal{L}(\theta^{(k)}), \theta \rangle + \frac{1}{2\eta_k} \|\theta - \theta^{(k)}\|_B^2 \right\} \\ &= \theta^{(k)} - \eta_k B^{-1} \nabla \mathcal{L}(\theta^{(k)})\end{aligned}$$

- ▶ use previous gradients to estimate B as diagonal matrix

$$G^{(k)} = \left(\sum_{i=1}^k \mathbf{diag} \left(\nabla \mathcal{L}(\theta^{(i)}) \right)^2 + \epsilon I \right)^{1/2}$$

AdaGrad (Duchi et al., 2011)

- view gradient descent as a first-order proximal method

$$\begin{aligned}\theta^{(k+1)} &= \operatorname{argmin}_{\theta} \left\{ \langle \nabla \mathcal{L}(\theta^{(k)}), \theta \rangle + \frac{1}{2\eta_k} \|\theta - \theta^{(k)}\|_2^2 \right\} \\ &= \theta^{(k)} - \eta_k \nabla \mathcal{L}(\theta^{(k)})\end{aligned}$$

- adapt the step size (geometry) for different features to increase influence of rare but informative features

$$\begin{aligned}\theta^{(k+1)} &= \operatorname{argmin}_{\theta} \left\{ \langle \nabla \mathcal{L}(\theta^{(k)}), \theta \rangle + \frac{1}{2\eta_k} \|\theta - \theta^{(k)}\|_B^2 \right\} \\ &= \theta^{(k)} - \eta_k B^{-1} \nabla \mathcal{L}(\theta^{(k)})\end{aligned}$$

- use previous gradients to estimate B as diagonal matrix

$$G^{(k)} = \left(\sum_{i=1}^k \mathbf{diag} \left(\nabla \mathcal{L}(\theta^{(i)}) \right)^2 + \epsilon I \right)^{1/2}$$

- weakness: $G^{(k)}$ keeps growing so learning rate becomes infinitesimally small

Adam

(Kingma and Ba, 2014)

- maintains exponentially decaying averages of past gradients and gradients-squared

$$m^{(k)} = \beta_1 m^{(k-1)} + (1 - \beta_1) \nabla \mathcal{L}(\theta^{(k)})$$

$$v^{(k)} = \beta_2 v^{(k-1)} + (1 - \beta_2) \nabla \mathcal{L}(\theta^{(k)})^2$$

Adam

(Kingma and Ba, 2014)

- maintains exponentially decaying averages of past gradients and gradients-squared

$$m^{(k)} = \beta_1 m^{(k-1)} + (1 - \beta_1) \nabla \mathcal{L}(\theta^{(k)})$$

$$v^{(k)} = \beta_2 v^{(k-1)} + (1 - \beta_2) \nabla \mathcal{L}(\theta^{(k)})^2$$

- compensate for bias (c.f. unbiased variance calculations)

$$\hat{m}^{(k)} = \frac{1}{1 - \beta_1^k} m^{(k)} \text{ and } \hat{v}^{(k)} = \frac{1}{1 - \beta_2^k} v^{(k)}$$

Adam

(Kingma and Ba, 2014)

- maintains exponentially decaying averages of past gradients and gradients-squared

$$m^{(k)} = \beta_1 m^{(k-1)} + (1 - \beta_1) \nabla \mathcal{L}(\theta^{(k)})$$

$$v^{(k)} = \beta_2 v^{(k-1)} + (1 - \beta_2) \nabla \mathcal{L}(\theta^{(k)})^2$$

- compensate for bias (c.f. unbiased variance calculations)

$$\hat{m}^{(k)} = \frac{1}{1 - \beta_1^k} m^{(k)} \text{ and } \hat{v}^{(k)} = \frac{1}{1 - \beta_2^k} v^{(k)}$$

- apply update rule

$$\theta^{(k+1)} = \theta^{(k)} - \eta \mathbf{diag} \left(\hat{v}^{(k)} + \epsilon \right)^{-1/2} \hat{m}^{(k)}$$

AdamW

(Loshchilov and Hutter, 2019)


- ▶ many other variants of update rules
- ▶ for example, Adam with weight decay

$$\theta^{(k+1)} = \theta^{(k)} - \eta \mathbf{diag} \left(\hat{v}^{(k)} + \epsilon \right)^{-1/2} \left(\hat{m}^{(k)} + w\theta^{(k)} \right)$$

- ▶ empirically works better than applying ℓ_2 regularization on $m^{(k)}$

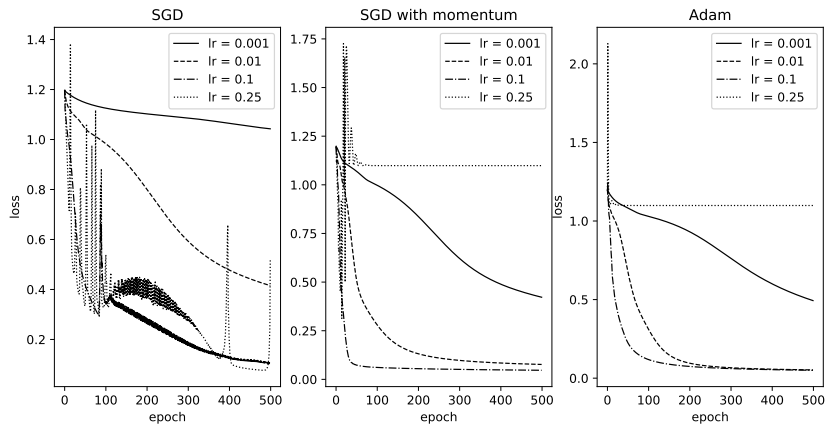
Learning Curves

- ▶ Iris dataset (3 classes, 4 features, 150 examples)
- ▶ multi-layer perceptron with:
 - ▶ 4 inputs, 8 hidden nodes (ReLU), 3 outputs (softmax)
- ▶ trained with cross-entropy loss (on all 150 examples)
- ▶ training loss plotted for different SGD variants



Learning Curves 2

- ▶ same setup as previous
- ▶ training loss plotted for different learning rates



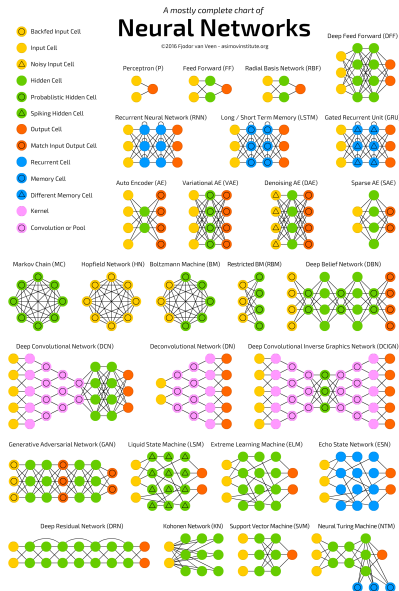
Training Tricks

Researchers have experimented with a number of tricks to help improve training/generalizability of deep models. Some of the tricks that seem to help include:

- ▶ careful parameter initialization (e.g., Xavier initialization)
- ▶ pre-training on a large annotated dataset (e.g., ImageNet)
- ▶ diverse mini-batches (e.g., randomize at start of epoch)
- ▶ data augmentation
- ▶ drop-out (not so popular anymore)
- ▶ batch normalization between layers
- ▶ gradient clipping
- ▶ intermediate supervision
- ▶ iterative refinement (skip/residual/short-cut connections)
- ▶ contrastive learning and multi-modal embedding

Model Zoo

(<http://www.asimovinstitute.org/neural-network-zoo/>)




Parameter Sharing

- ▶ consider a MLP acting on input $x = (x_1, x_2) \in \mathbb{R}^{2n}$

$$z = f \left(\begin{bmatrix} a_1^T & 0 \\ 0 & a_2^T \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} + \begin{bmatrix} b_1 \\ b_2 \end{bmatrix} \right) \in \mathbb{R}^2$$


$$y = f(c^T z + d) \in \mathbb{R}$$



- ▶ now suppose we want to share weights $a_1 = a_2 = a \in \mathbb{R}^n$
- ▶ then we need to combine gradients

$$\nabla_a y = \frac{\partial y}{\partial z_1} \nabla_a z_1 + \frac{\partial y}{\partial z_2} \nabla_a z_2$$

Parameter Sharing — Convolutions




- ▶ parameters vector $a \in \mathbb{R}^h$ and input $x \in \mathbb{R}^n$
- ▶ weighted sum of inputs becomes convolution with a shared filter kernel

$$z = f(a * x + b)$$

where $z_i = f(a^T x_{i+h:i} + b)$ for $i = 1, \dots, n - h$

Convolutional Neural Networks (CNNs)

- ▶ a CNN takes as input an image ($3 \times W \times H$ tensor) and performs successive layers of:
 - ▶ convolution, $(x * a)_{st} = \sum_{i,j} x_{s-i,t-j} \cdot a_{ij} + b$
 - ▶ non-linear transform, e.g., ReLU
 - ▶ pooling/downsampling, e.g., $z_{st} = \max\{x_{ij} \mid i, j \in \mathcal{N}_{st}\}$
- ▶ the final layers are typically “fully-connected” (i.e., an MLP)



- ▶ size of a is often called the “receptive field”
- ▶ variants: change stride/padding, dilated (atrous), decomposed, etc.

Recurrent Neural Networks

- ▶ we can feed the (partial) output of a network back into itself and time-step inference to give a recurrent network,

$$(y_t, h_t) = f(x_t, h_{t-1})$$

- ▶ useful for sequence modelling

Recurrent Neural Networks

- ▶ we can feed the (partial) output of a network back into itself and time-step inference to give a recurrent network,

$$(y_t, h_t) = f(x_t, h_{t-1})$$

- ▶ useful for sequence modelling
- ▶ most popular variant is the *long short-term memory* (LSTM) network

$$f_t = \sigma(W_f x_t + U_f h_{t-1} + b_f)$$

forget activations

$$i_t = \sigma(W_i x_t + U_i h_{t-1} + b_i)$$

input activations

$$o_t = \sigma(W_o x_t + U_o h_{t-1} + b_o)$$

output activations

$$c_t = f_t \circ c_{t-1} + i_t \circ \sigma(W_c x_t + U_c h_{t-1} + b_c)$$

cell memory

$$h_t = o_t \circ \sigma(c_t)$$

cell output

Back-propagation Through Time

- We can train recurrent networks by “unrolling” the network to a given time horizon and back-propagating through time.

$$\begin{aligned}(y_t, h_t) &= f(x_t, h_{t-1}; \theta) \\ (y_{t-1}, h_{t-1}) &= f(x_{t-1}, h_{t-2}; \theta) \\ &\vdots \\ (y_1, h_1) &= f(x_1, h_0; \theta)\end{aligned}$$

Back-propagation Through Time

- ▶ We can train recurrent networks by “unrolling” the network to a given time horizon and back-propagating through time.


$$\begin{aligned}(y_t, h_t) &= f(x_t, h_{t-1}; \theta) \\ (y_{t-1}, h_{t-1}) &= f(x_{t-1}, h_{t-2}; \theta) \\ &\vdots \\ (y_1, h_1) &= f(x_1, h_0; \theta)\end{aligned}$$

- ▶ Then (assuming the loss is applied to the last output, y_t)

$$\begin{aligned}\frac{\partial \mathcal{L}}{\partial \theta} &= \frac{\partial \mathcal{L}}{\partial y_t} \frac{\partial y_t}{\partial \theta} + \frac{\partial \mathcal{L}}{\partial y_t} \frac{\partial y_t}{\partial h_{t-1}} \frac{\partial h_{t-1}}{\partial \theta} + \dots \\ &\quad \dots + \frac{\partial \mathcal{L}}{\partial y_t} \frac{\partial y_t}{\partial h_{t-1}} \left(\prod_{\tau=1}^{t-1} \frac{\partial h_{\tau+1}}{\partial h_{\tau}} \right) \frac{\partial h_1}{\partial \theta}\end{aligned}$$

Short-cut Connections

- ▶ one way to address the issue of diminishing gradients is via short-cut connections
- ▶ the ResNet block is a classic example




- ▶ the gradient through the ResNet block is



$$\frac{\partial y}{\partial x} = \mathcal{F}'(x) + 1$$

- ▶ ResNet also models “residual signals”, which may help lead to its empirically observed improved performance



Machine Learning: An Abstract View



Machine Learning: An Abstract View




Machine Learning: Encoder-decoder Models



Auto-encoders

- ▶ attempts to learn a low-dimensional (and sometimes sparse) representation of data
- ▶ does not need explicit supervision



- ▶ encoder: $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$ with $m \ll n$
- ▶ decoder: $g : \mathbb{R}^m \rightarrow \mathbb{R}^n$
- ▶ learn such that $g(f(x)) \approx x$
- ▶ then $z = f(x)$ is a good latent representation of x

Encoder-decoder Models

- ▶ we can generalize auto-encoders where the output does not need to be a reconstruction of the input
- ▶ still learned end-to-end so that the “latent” representation captures some higher level meaning
- ▶ this gives rise to many interesting applications, e.g.,
 - ▶ use a CNN encoder-decoder for semantic segmentation or style-transfer
 - ▶ use an RNN encoder-decoder for language translation (sometimes called sequence-to-sequence models)
 - ▶ use a mixed CNN-RNN encoder-decoder for image captioning
 - ▶ learn joint image-language embedding (e.g., CLIP) then generate images from natural language description (DALL-E, Imagen, Stable Diffusion)

Probabilistic Auto-encoders

- ▶ adds probability distributions over the data x and latent representation z , $p(x | z)$ and $p(z)$, respectively
- ▶ some things we might want to are:
 - ▶ jointly sample data and representation from $p_\theta(x, z)$
 - ▶ compute the marginal $p_\theta(x) = \int p_\theta(x | z)p_\theta(z)dz$
 - ▶ sample representation from $p_\theta(z | x) = \frac{p_\theta(x|z)p_\theta(z)}{p_\theta(x)}$
 - ▶ infer (learn) parameters θ
- ▶ the first is easy, the last three are difficult (in general)
- ▶ solved by learning a distribution $q_\phi(z | x)$

Variational Auto-encoders (VAEs)

(Kingma and Welling, 2013; Rezende et al., 2014)

- ▶ admits efficient sampling from a learned distribution, $p_{\theta}(x | z)$,

$$z \sim \mathcal{N}(0, I)$$

$$x \sim \mathcal{N}(\mu(z), \Sigma(z))$$

where functions $\mu(\cdot)$ and $\Sigma(\cdot)$ are learned (decoder models)

- ▶ introduce variational lower bound for jointly learning ϕ and θ

$$\log p_{\theta}(x) \geq \mathbf{E}_{q_{\phi}(z|x)} \left[\log \frac{p_{\theta}(x, z)}{q_{\phi}(z | x)} \right]$$

- ▶ the gradients cannot propagate through random variables¹ so we need to reparametrize x as follows

$$x \sim \mathcal{N}(\mu, \Sigma) \text{ becomes } x = \mu + L\epsilon$$

where $\epsilon \sim \mathcal{N}(0, I)$ and $\Sigma = LL^T$

¹sweeping some technicalities under the carpet

Generative Adversarial Models (GANs)

(Goodfellow et al., 2014)

- ▶ alternative method for learning to generate samples based on ideas of training competing models to reach an equilibrium
- ▶ does not explicitly represent the data generating distribution
- ▶ the **generator** G_θ produces novel “realistic” samples that fool the discriminator
- ▶ the **discriminator** D_ϕ classifies samples as “real” or “fake”
- ▶ mathematically we can write the objective

$$\min_{\theta} \max_{\phi} [E_{x \sim p_{\text{data}}} \log D_{\phi}(x) + E_{z \sim p_Z} \log(1 - D_{\phi}(G_{\theta}(z)))]$$

- ▶ due to optimization issues, we typically replace generative loss with $-\log D_{\phi}(G_{\theta}(z))$

Attention and Transformers

(Polosukhin et al., 2017)

- ▶ recurrent models on sequences tend to have difficulties learning due to vanishing/exploding gradients (LSTMs alleviate this to some extent)
- ▶ solution is to replace recurrent network with **attention** mechanism

$$f(Q, K, V) = \text{softmax} \left(\frac{QK^T}{\sqrt{n}} \right) V$$

where $Q \in \mathbb{R}^{m \times n}$, $K \in \mathbb{R}^{p \times n}$ and $V \in \mathbb{R}^{p \times q}$ are the “query”, “key” and “value” matrices, respectively.

- ▶ add positional encoding (multi-frequency sine waves) for longer sequences
- ▶ (add bells and whistles for full transformer model)

Diffusion Models

(Ho, Jain and Abbeel, 2020)

► forward process:

$$x_0 \sim q(x) \rightarrow \cdots \rightarrow x_{t-1} \xrightarrow{q(x_t|x_{t-1})} x_t \rightarrow \cdots \rightarrow x_T$$

where $q(x_t | x_{t-1}) = \mathcal{N}(\sqrt{1 - \beta_{t-1}}x_{t-1}, \beta_{t-1}I)$

► reverse process:

$$x_t \sim \mathcal{N}(0, I) \rightarrow \cdots \rightarrow x_t \xrightarrow{q(x_{t-1}|x_t)} x_{t-1} \rightarrow \cdots \rightarrow x_0$$

but where $q(x_{t-1} | x_t)$ is, in general, unknown

- uses variational l.b. to estimate $q(x_{t-1} | x_t)$ as $p_\theta(x_{t-1} | x_t)$
- here p_θ is a learned denoising model
- can be used for image generation (or guided image generation with a few modifications to p_θ)

Reinforcement Learning

- ▶ at each timestep an agent observes the state of the environment s_t and takes an action a_t ; the environment provides a reward r_{t+1} and updates its state s_{t+1}
- ▶ r_{t+1} and s_{t+1} may be governed by stochastic processes
- ▶ reinforcement learning results in a policy π that aims to maximize the sum of discounted future rewards

$$R = \sum_{t=0}^{\infty} \gamma^t r_t$$

- ▶ we often learn a Q -function from experience

$$Q^{\pi}(s, a) = \mathbf{E}[R \mid s, a, \pi]$$

which gives the expected reward R for taking action a in state s and then following policy π

- ▶ we obtain the optimal action as $a_t \in \operatorname{argmax} Q^{\pi^*}(s_t, \cdot)$

Deep Reinforcement Learning

(Mni et al., Nature, 2015)

deep Q-learning (DQN):

- ▶ approximate Q^{π^*} by a neural network f_{θ}
- ▶ stochastic gradient update of parameters using loss

$$\mathcal{L}(\theta) = \left(\overbrace{r + \gamma \max_{a'} f_{\tilde{\theta}}(s', a')}^{\text{one action on } \tilde{Q}} - \underbrace{f_{\theta}(s, a)}_{\text{learned } Q} \right)^2$$

where $\tilde{\theta}$ are the fixed previous model parameters and a, s, s' and r are sampled action, state, next state, and final reward from experience, respectively

policy gradient: directly learn $a_t = f(x_t; \theta)$

- ▶ e.g., $a_t \sim p_{\theta}(\text{"win"} \mid a_t, s_t) \propto p_{\theta}(a_t \mid s_t, \text{"win"})$
- ▶ train by rolling out sampled games and weighting each action by the final total reward